Agent BizOps:

Agentifying AI Adoption with Enterprise Digital Twin

Version 0.5.0

Published by Tavro AI, LLC

July 27, 2026

Table of Contents

License and Usage

The information provided in this document does not, and is not intended to, constitute legal advice. All information is for general informational purposes only. This document may contain links to third-party websites provided solely for convenience. Tavro AI, LLC ("Tavro") does not recommend or endorse the contents of these third-party sites.

Acknowledgements

Leader Authors:

- Sunil Soares (Tavro AI)
- Sanjeev Varma (Tavro AI)

Key Contributors and Reviewers

The following individuals are acknowledged for their significant contributions to the review and development of this document, **listed alphabetically by last name**:

Key Contributor / Reviewer Name	Affiliation
Prasanna Kumar Akiri	Tavro AI
Raj Arumugam	Entergy
Venkata Atluri	USAA
Ashutosh Atre	Independent Contributor
Appal Badi	Independent Contributor
Val Calvo	AgFirst Bank
Madhuri Challapalli	Independent Contributor
Stan Christiaens	Collibra
Joe DeLuca	Independent Contributor
Karan Dhawal	ZS Associates
Yoshita Dhillon	Independent Contributor
Tony DiPerna	BankUnited
Mihir Dudhatra	Tavro AI
Kjetil Eritzland	Capgemini Norway
Abeda Essop	Sanlam South Africa
Kevin Ferry	Natixis Investment Advisors
Forrest Gilman	Independent Contributor
Pierre Gomes	BankUnited
Vikas Gupta	Hindustan Petroleum

Shawn Harbaugh	Citizens & Northern Bank
Ravindra Harve	Boston College
Mike Jennings	Independent Contributor (Formerly Walgreens)
Peter Kapur	CarMax
Michael Koegler	Market Alpha Advisors
Aditya Kongara	Independent Contributor
Manirajkumar Kotha	Tavro AI
Mmatseleng Lefike	Sanlam South Africa
Jiten Mehta	Independent Contributor
Gokula Mishra	OmniProAI (Formerly McDonald’s)
Fahad Moosa	Independent Contributor
Rahul Pandit	Tavro AI
Arpan Patel	Delta Community Credit Union
Nick Raad	Independent Contributor
Su Rayburn	Delta Community Credit Union
Shyam Rasaily	T. Rowe Price
Phani Rayudu	Walgreens
Philip Reed	Independent Contributor
Khushboo Shah	Tavro AI
Doug Shannon	Independent Contributor
Jen Shomshor	Independent Contributor
Marti Smith	Independent Contributor
Liz Soares	Independent Contributor
Bryan Swann	Independent Contributor
Liv Watson	Independent Contributor
John Yelle	Independent Contributor (Formerly DTCC)

What is Agent BizOps?

Every enterprise wants to adopt AI agents, but few have the operational discipline to do it well.

Agent BizOps is the practice of agentifying AI adoption using an Enterprise Digital Twin.

Agent BizOps gives business and technology leaders a shared operating model so that AI initiatives move from idea to production with full visibility into cost, risk, and business impact. According to McKinsey, 62 percent of organizations are in the early stages of AI adoption.[1] A root cause of this slow progress is the absence of a common language across business, technology, risk, and finance teams. Without that shared context, agent projects stall in pilot mode, duplicated efforts multiply, and governance becomes an afterthought. This is precisely the gap an Enterprise Digital Twin fills, and it is why Agent BizOps is important.

Agent adoption is an uphill battle

---

1 Quantum Black AI by McKinsey, “The State of AI in 2025: Agents, innovation, and transformation,” https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai.

What is an Enterprise Digital Twin?

An **Enterprise Digital Twin** is a living metadata model of your organization. It connects the information that different teams already maintain — strategy documents, application inventories, process maps, financial reports, risk registers, master data — into a single, queryable layer. Rather than scattering this knowledge across spreadsheets and wikis, the Enterprise Digital Twin makes it available as context for AI agent planning, development, and governance.

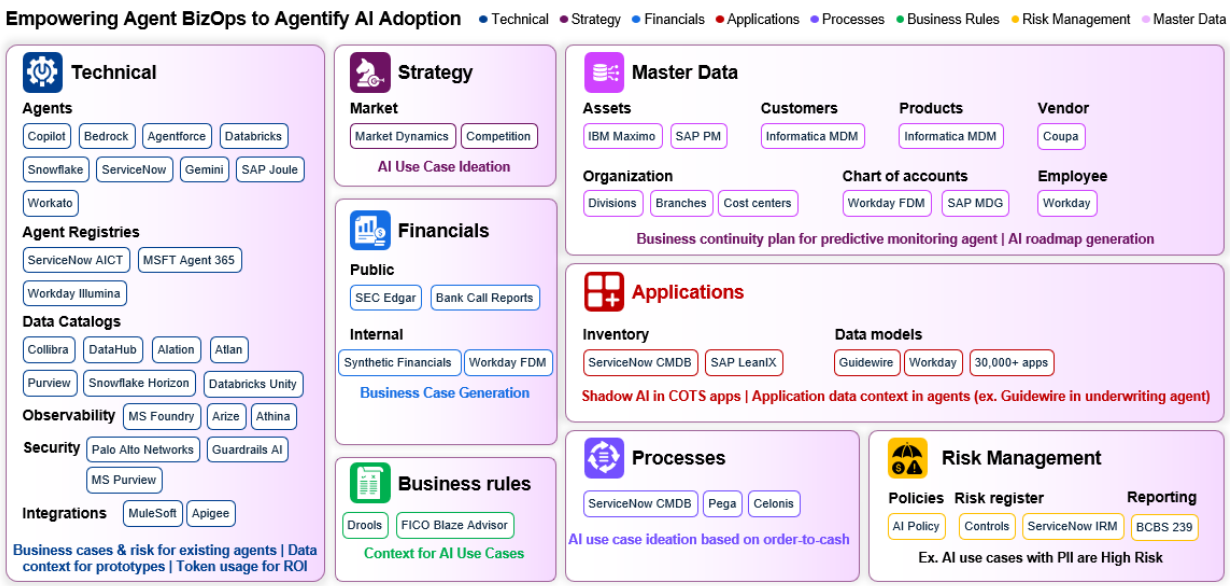

Overview of Enterprise Digital Twin

The Enterprise Digital Twin consists of the following information:

Category	Sub-Category	Description
Technical Information	Agents	Metadata on agents (e.g., Microsoft Copilot, Amazon Bedrock, Salesforce Agentforce, Databricks, Snowflake Cortex, ServiceNow NowAssist, SAP Joule, Workato)
	Agent Registries	Metadata from agent registries (e.g., ServiceNow AI Control Tower, Microsoft Agent 365, Workday Illumina), related to the agent sub-category above
	Data Catalogs	Metadata from data catalogs (e.g., Collibra, DataHub, Alation, Atlan, Microsoft Purview, Snowflake Horizon Data Catalog, Databricks Unity Catalog)
	Observability	Token usage and other session metadata from observability tools (e.g., Microsoft Foundry, Arize, Athina)
	Security	Security-related signals from platforms like Palo Alto Networks, Guardrails AI, Microsoft Purview, etc.

	Integrations	Information from integration platforms like Apigee, MuleSoft, etc.
Strategy	Market	Information like market dynamics, competitive intelligence, etc.
Financials	Public	Publicly-available financials (e.g., SEC Edgar filings like 10K, 10Q, annual reports, bank call reports, etc.)
	Internal	Internal reports (e.g., board presentations, Workday FDM reports, etc.)
Business Rules	Inventory	Business rules within PDFs or platforms like Drools and FICO Blaze Advisor
Master Data	Assets	Asset master data within platforms like IBM Maximo and SAP PM
	Customers	Customer master data within platforms like Informatica MDM
	Products	Product master data within platforms like Informatica MDM
	Vendors	Vendor master data within platforms like Coupa
	Organization	Organizational master data (e.g., divisions, branches, cost centers, etc.)
	Chart of Accounts	Chart of accounts within platforms like Workday FDM and SAP MDG
	Employee	Employee master data within platforms like Workday
Applications	Inventory	Application inventory within platforms like ServiceNow CMDB and SAP LeanIX
	Data Models	Data models for thousands of COTS and RYO applications (e.g., Guidewire, Workday)
Processes	Inventory	Process inventory within platforms like ServiceNow CMDB, Pega, Celonis, etc.
Risk Management	Policies	Risk management policies (e.g., AI policy, data policy, etc.)
	Risk Register	Artifacts such as risks and controls within GRC platforms like MetricStream, ServiceNow IRM, etc.
	Reporting	Risk and regulatory reports such as those related to BCBS 239

Components of Enterprise Digital Twin

Do We Need to Build Out the Enterprise Digital Twin All At Once?

Because the Enterprise Digital Twin may initially seem daunting, it leads to the logical question, "Do we need to build it all out at once?" The answer is a resounding no for the following reasons:

1. **Enterprises have been building "digital twins" for decades**. This has occurred as part of process automation so the concept of an Enterprise Digital Twin is not new.

2. **AI makes the digital twin more active, rather than passive.** You can still do the same task, may be even more cost effectively, but now this digital twin can answer questions like "what's going on?," and "when will I get what I want?"

   For example, the image below shows an Enterprise Digital Twin for a 16-physician urology practice that is independently owned and operated. The Enterprise Digital Twin includes the company profile, strategy, organization, financials, risks, applications, and processes.

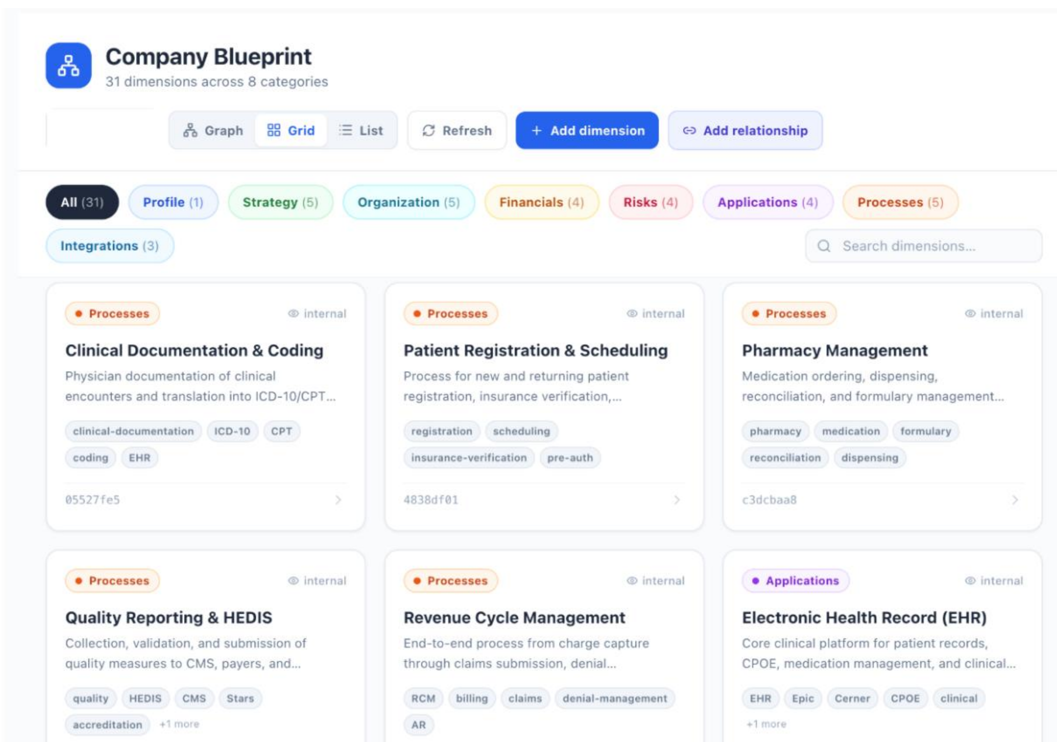

Enterprise Digital Twin at urology practice

The Enterprise Digital Twin provides the context so that AI can spark ideas for AI use cases that *"provide quick return but do not involve a huge upfront investment."*

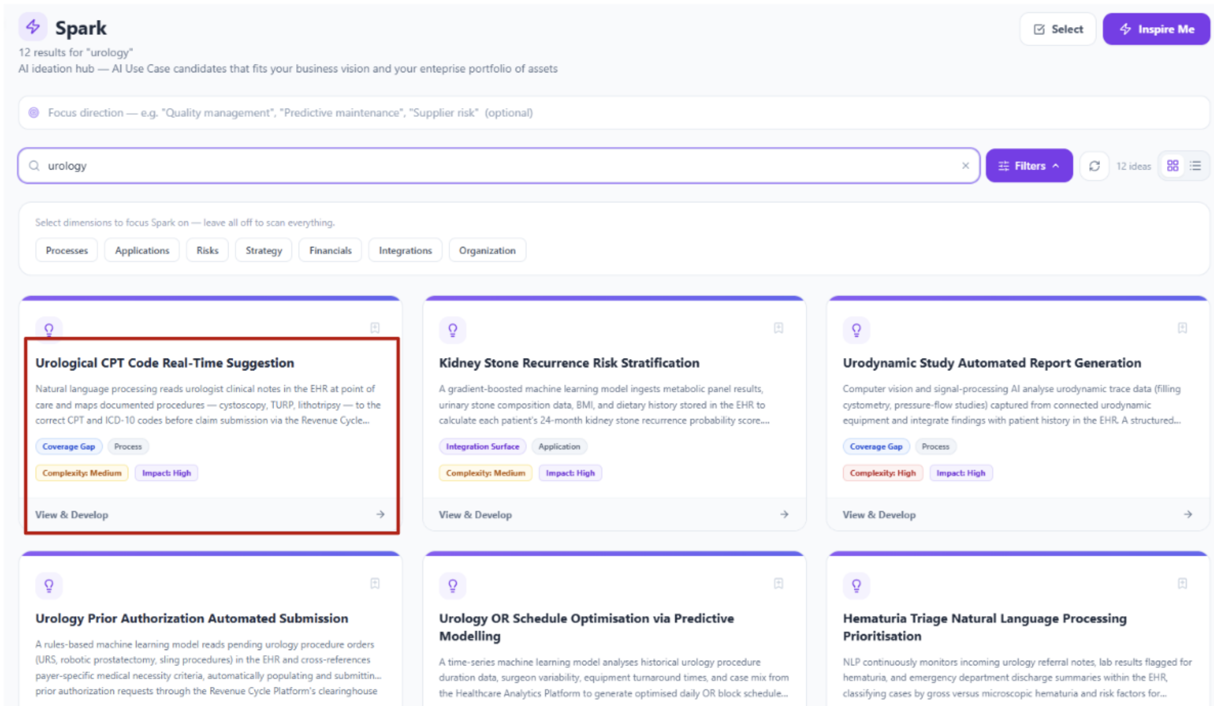

Ideation of AI use cases based on Enterprise Digital Twin at urology practice

3. **Constant enrichment results in improved outcomes.** For example, adding financials to the Enterprise Digital Twin enriches the business case for **Regulatory Change Classification and Policy Mapping** AI use case at a community bank.

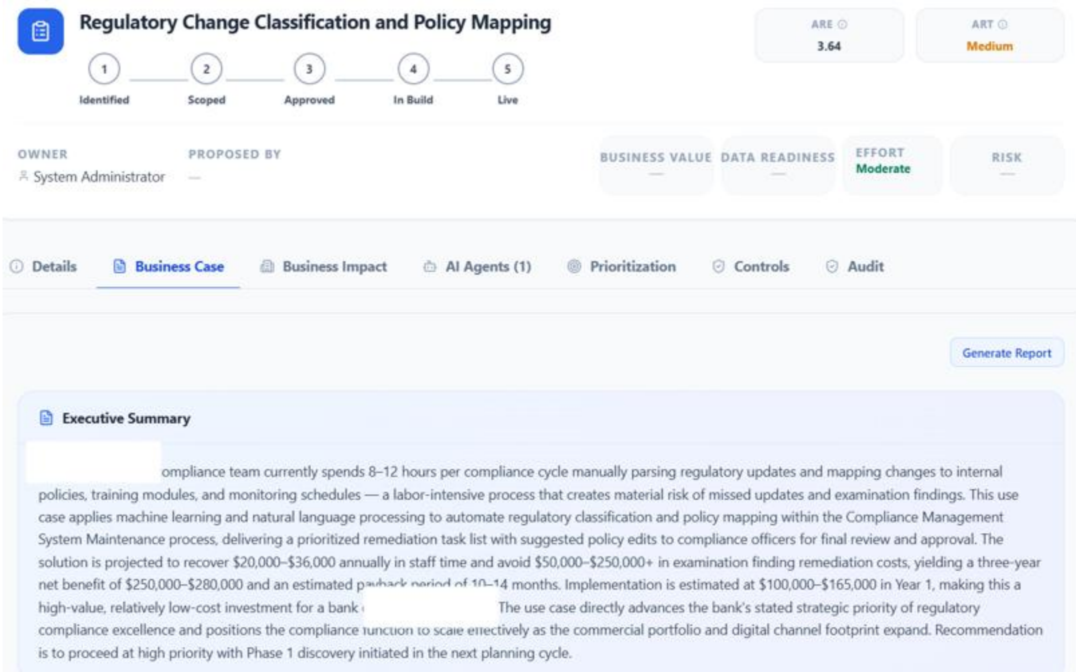

AI use case at community bank includes financial benefits based on enriched Enterprise Digital Twin

4. **Cross-organization efficiencies and improvements accrue when multiple functions enrich the Enterprise Digital Twin.** For example, the **Regulatory Change Classifier and Policy Agent** implements the AI use case at the abovementioned community bank. The **Finance** team drove the quantification of the financial benefits of the AI use case as discussed in step 3 above. The **Risk Management** team scored the agent as **3.64 – Medium Risk**. This supports cross-organizational trade-offs.

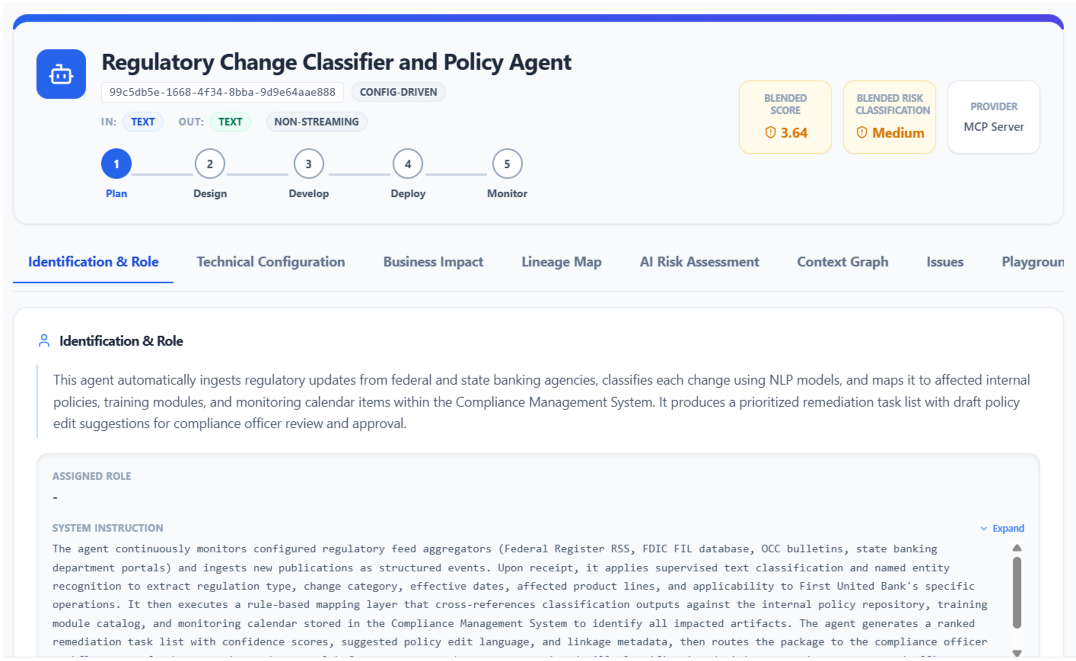

Risk assessment for AI agent at community bank

5. **Company leadership can answer strategic questions when every department has enriched the Enterprise Digital Twin in step 3 above.** These questions may relate to company strategy, restructuring, prioritization of AI uses, and the AI roadmap. For example, AI use cases at a community bank may be automatically prioritized based on effort, risk, and ROI based on the Enterprise Digital Twin.

Five AI use cases evaluated across Effort, Risk, and ROI. One use case (Commercial Risk Rating Migration Scoring) scores as a Money pit algorithmically but is recommended for approval as a strategic override.

#	Use case	Effort	Risk	ROI	Quadrant
1	Regulatory Change Classification and Policy Mapping	2	3	3	**Quick win**
2	Retail Deposit Churn Prediction and Retention Scoring	2	4	3	**Fill-in**
3	Currency Transaction Report (CTR) Aggregation Anomaly Detection	3	5	3	**Fill-in**
4	CECL Loss Rate Forecasting with Machine Learning	3	4	2	**Fill-in**
5	Predictive Risk Rating Migration Scoring (Commercial Portfolio)	4	4	5	**Override (!)**

(!) A note on overrides: Predictive Risk Rating Migration Scoring for the Commercial Portfolio carries the highest implementation cost and Effort/Risk profile in the set (algorithmically a Money pit), but its projected 3-year net benefit of $4.3M–$8.7M and 215%–435% ROI — the strongest financial case of the five — support approval as a strategic override, proceeding in phases with formal governance sign-off.

Prioritization of AI use cases at a community bank based on Enterprise Digital Twin

The table below shows how various categories of the Enterprise Digital Twin may be used to drive business value.

Category	Usage Examples
Technical Information	• Data catalogs provide the data context for AI use cases • Observability platform provide information on token usage that informs ROI for AI use cases
Strategy	• Enterprise strategy informs the ideation of AI use cases
Financials	• Enterprise financials should drive the business cases for AI
Business Rules	• Business rules may provide context for AI use cases (e.g., a business rule that high-risk AI use cases must be approved by the AI Steering Committee)
Master Data	• For example, a utility has a predictive monitoring AI use case with mappings of the agents to associated applications and power transmission assets. This allows the utility to quickly generate a business continuity plan for the predictive monitoring agent.
Applications	• For example, an insurance carrier is developing underwriting agents that

	use Guidewire Policy Center. This allows the utility to quick determine the readiness of Guidewire data to support the AI use case.
Processes	• For example, an order-to-cash process may support the ideation of multiple AI use cases
Risk Management	• For example, the enterprise's risk appetite may determine that AI use cases with personally identifiable information are automatically classified as high risk

Importance of the Enterprise Digital Twin

Who cares about the Enterprise Digital Twin?

The Enterprise Digital Twin is relevant to every executive who touches AI adoption.

CEOs and other executives with P&L responsibility prioritize AI use cases and roadmaps. CFOs need it to justify AI investments with hard ROI data. Chief AI Officers use it to support an AI agent factory.

Control functions like the Chief Risk Officer and Data & AI Governance Leads use it to track agent sprawl and infrastructure costs.

And line-of-business leaders use it to identify where agents can have the highest impact on their operations. The image below shows how different stakeholders engage with the Enterprise Digital Twin.

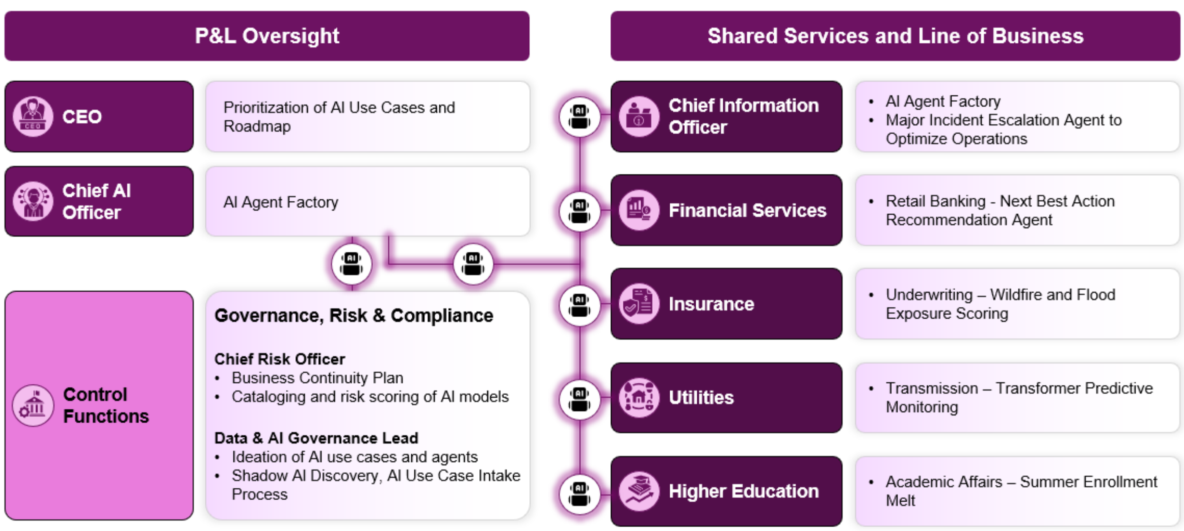

Stakeholders for the Enterprise Digital Twin

Evolution of Agent BizOps Standard

The Agent BizOps Standard started out as the **Agent Metadata Specification**. Based on feedback from over a hundred industry professionals across multiple industries, the focus of the standard evolved from agent metadata to the use of metadata (Enterprise Digital Twin) to support AI adoption.